




Boosting

Lecturer: Sang Hwa Lee



What is AdaBoost ?

❖ Concept of AdaBoosting

AdaBoost is an algorithm for constructing a “strong” classifier as linear combination of “simple” “weak” classifiers $h_t(x): X \rightarrow \{-1, +1\}$.

$$f(x) = \sum_{t=1}^T \alpha_t h_t(x)$$


❖ Terminology

$h_t(x)$... “weak” or basis classifier, hypothesis, “feature”

$H(x) = \text{sign}(f(x))$... “strong” or final classifier/hypothesis

AdaBoost Algorithm

Given: $(x_1, y_1), \dots, (x_m, y_m); x_i \in X, y_i \in \{-1, +1\}$




AdaBoost Algorithm

Given: $(x_1, y_1), \dots, (x_m, y_m); x_i \in X, y_i \in \{-1, +1\}$

Initialise weights $D_1(i) = 1/m$

For $t = 1, \dots, T$:




AdaBoost Algorithm

Given: $(x_1, y_1), \dots, (x_m, y_m); x_i \in X, y_i \in \{-1, +1\}$

Initialise weights $D_1(i) = 1/m$

For $t = 1, \dots, T$:

◇ Find $h_t = \arg \min_{h_j \in \mathcal{H}} \epsilon_j = \sum_{i=1}^m D_t(i) \mathbb{I}[y_i \neq h_j(x_i)]$



AdaBoost Algorithm


Given: $(x_1, y_1), \dots, (x_m, y_m); x_i \in X, y_i \in \{-1, +1\}$

Initialise weights $D_1(i) = 1/m$

For $t = 1, \dots, T$:

- ◇ Find $h_t = \arg \min_{h_j \in \mathcal{H}} \epsilon_j = \sum_{i=1}^m D_t(i) \mathbb{I}[y_i \neq h_j(x_i)]$
- ◇ If $\epsilon_t \geq 1/2$ then stop

$t = 1$



AdaBoost Algorithm


Given: $(x_1, y_1), \dots, (x_m, y_m); x_i \in X, y_i \in \{-1, +1\}$

Initialise weights $D_1(i) = 1/m$

For $t = 1, \dots, T$:

- ◇ Find $h_t = \arg \min_{h_j \in \mathcal{H}} \epsilon_j = \sum_{i=1}^m D_t(i) \mathbb{I}[y_i \neq h_j(x_i)]$
- ◇ If $\epsilon_t \geq 1/2$ then stop
- ◇ Set $\alpha_t = \frac{1}{2} \log\left(\frac{1-\epsilon_t}{\epsilon_t}\right)$

$t = 1$



AdaBoost Algorithm

Given: $(x_1, y_1), \dots, (x_m, y_m); x_i \in X, y_i \in \{-1, +1\}$

Initialise weights $D_1(i) = 1/m$


For $t = 1, \dots, T$:

- ◇ Find $h_t = \arg \min_{h_j \in \mathcal{H}} \epsilon_j = \sum_{i=1}^m D_t(i) \mathbb{I}[y_i \neq h_j(x_i)]$
- ◇ If $\epsilon_t \geq 1/2$ then stop
- ◇ Set $\alpha_t = \frac{1}{2} \log\left(\frac{1-\epsilon_t}{\epsilon_t}\right)$
- ◇ Update

$$D_{t+1}(i) = \frac{D_t(i) \exp(-\alpha_t y_i h_t(x_i))}{Z_t}$$

where Z_t is normalisation factor

$t = 1$



AdaBoost Algorithm

Given: $(x_1, y_1), \dots, (x_m, y_m); x_i \in X, y_i \in \{-1, +1\}$

Initialise weights $D_1(i) = 1/m$

For $t = 1, \dots, T$:

◇ Find $h_t = \arg \min_{h_j \in \mathcal{H}} \epsilon_j = \sum_{i=1}^m D_t(i) \mathbb{I}[y_i \neq h_j(x_i)]$

◇ If $\epsilon_t \geq 1/2$ then stop

◇ Set $\alpha_t = \frac{1}{2} \log\left(\frac{1-\epsilon_t}{\epsilon_t}\right)$

◇ Update



$$D_{t+1}(i) = \frac{D_t(i) \exp(-\alpha_t y_i h_t(x_i))}{Z_t}$$

where Z_t is normalisation factor

Output the final classifier:

$$H(x) = \text{sign} \left(\sum_{t=1}^T \alpha_t h_t(x) \right)$$

$t = 1$



AdaBoost Algorithm

Given: $(x_1, y_1), \dots, (x_m, y_m); x_i \in X, y_i \in \{-1, +1\}$

Initialise weights $D_1(i) = 1/m$

For $t = 1, \dots, T$:

◇ Find $h_t = \arg \min_{h_j \in \mathcal{H}} \epsilon_j = \sum_{i=1}^m D_t(i) \mathbb{I}[y_i \neq h_j(x_i)]$

◇ If $\epsilon_t \geq 1/2$ then stop

◇ Set $\alpha_t = \frac{1}{2} \log\left(\frac{1-\epsilon_t}{\epsilon_t}\right)$

◇ Update



$$D_{t+1}(i) = \frac{D_t(i) \exp(-\alpha_t y_i h_t(x_i))}{Z_t}$$

where Z_t is normalisation factor

Output the final classifier:

$$H(x) = \text{sign} \left(\sum_{t=1}^T \alpha_t h_t(x) \right)$$

$t = 2$



AdaBoost Algorithm

Given: $(x_1, y_1), \dots, (x_m, y_m); x_i \in X, y_i \in \{-1, +1\}$

Initialise weights $D_1(i) = 1/m$

For $t = 1, \dots, T$:

◇ Find $h_t = \arg \min_{h_j \in \mathcal{H}} \epsilon_j = \sum_{i=1}^m D_t(i) \mathbb{I}[y_i \neq h_j(x_i)]$

◇ If $\epsilon_t \geq 1/2$ then stop

◇ Set $\alpha_t = \frac{1}{2} \log\left(\frac{1-\epsilon_t}{\epsilon_t}\right)$

◇ Update



$$D_{t+1}(i) = \frac{D_t(i) \exp(-\alpha_t y_i h_t(x_i))}{Z_t}$$

where Z_t is normalisation factor

Output the final classifier:

$$H(x) = \text{sign} \left(\sum_{t=1}^T \alpha_t h_t(x) \right)$$

$t = 3$



AdaBoost Algorithm

Given: $(x_1, y_1), \dots, (x_m, y_m); x_i \in X, y_i \in \{-1, +1\}$

Initialise weights $D_1(i) = 1/m$

For $t = 1, \dots, T$:

◇ Find $h_t = \arg \min_{h_j \in \mathcal{H}} \epsilon_j = \sum_{i=1}^m D_t(i) \mathbb{I}[y_i \neq h_j(x_i)]$

◇ If $\epsilon_t \geq 1/2$ then stop

◇ Set $\alpha_t = \frac{1}{2} \log\left(\frac{1-\epsilon_t}{\epsilon_t}\right)$

◇ Update



$$D_{t+1}(i) = \frac{D_t(i) \exp(-\alpha_t y_i h_t(x_i))}{Z_t}$$

where Z_t is normalisation factor

Output the final classifier:

$$H(x) = \text{sign} \left(\sum_{t=1}^T \alpha_t h_t(x) \right)$$

$t = 6$



AdaBoost Algorithm

Given: $(x_1, y_1), \dots, (x_m, y_m); x_i \in X, y_i \in \{-1, +1\}$

Initialise weights $D_1(i) = 1/m$

For $t = 1, \dots, T$:

- ◇ Find $h_t = \arg \min_{h_j \in \mathcal{H}} \epsilon_j = \sum_{i=1}^m D_t(i) \mathbb{I}[y_i \neq h_j(x_i)]$
- ◇ If $\epsilon_t \geq 1/2$ then stop
- ◇ Set $\alpha_t = \frac{1}{2} \log\left(\frac{1-\epsilon_t}{\epsilon_t}\right)$
- ◇ Update



$$D_{t+1}(i) = \frac{D_t(i) \exp(-\alpha_t y_i h_t(x_i))}{Z_t}$$

where Z_t is normalisation factor

Output the final classifier:

$$H(x) = \text{sign} \left(\sum_{t=1}^T \alpha_t h_t(x) \right)$$

$t = 40$




Reweighting

Effect on the training set

$$D_{t+1}(i) = \frac{D_t(i) \exp(-a_t y_i h_t(x_i))}{Z_t}$$

$$\exp(-\alpha_t y_i h_t(x_i)) \begin{cases} < 1, & y_i = h_t(x_i) \\ > 1, & y_i \neq h_t(x_i) \end{cases}$$

- ⇒ Increase (decrease) weight of wrongly (correctly) classified examples
- ⇒ The weight is the upper bound on the error of a given example
- ⇒ All information about previously selected “features” is captured in D_t



Upper Bound Theorem (1)

Theorem: The following upper bound holds on the training error of H


$$\frac{1}{m} |\{i : H(x_i) \neq y_i\}| \leq \prod_{t=1}^T Z_t$$

Proof: By unravelling the update rule

$$\begin{aligned} D_{T+1}(i) &= \frac{D_t(i) \exp(-a_t y_i h_t(x_i))}{Z_t} \\ &= \frac{\exp(-\sum_t \alpha_t y_i h_t(x_i))}{m \prod_t Z_t} = \frac{\exp(-y_i f(x_i))}{m \prod_t Z_t} \end{aligned}$$

If $H(x_i) \neq y_i$ then $y_i f(x_i) \leq 0$ implying that $\exp(-y_i f(x_i)) \geq 1$, thus

$$\begin{aligned} \mathbb{I}[H(x_i) \neq y_i] &\leq \exp(-y_i f(x_i)) \\ \frac{1}{m} \sum_i \mathbb{I}[H(x_i) \neq y_i] &\leq \frac{1}{m} \sum_i \exp(-y_i f(x_i)) \\ &= \sum_i \left(\prod_t Z_t \right) D_{T+1}(i) = \prod_t Z_t \end{aligned}$$



Upper Bound Theorem (2)

- ❖ Instead of minimising the training error, its upper bound can be minimize
d. This can be done by minimising Z_t in each training round by:
 - Choosing optimal h_t , and
 - Finding optimal a_t
- ❖ AdaBoost can be proved to maximise margin
- ❖ AdaBoost iteratively fits an additive logistic regression model

Choosing optimal α_t

We attempt to minimise $Z_t = \sum_i D_t(i) \exp(-\alpha_t y_i h_t(x_i))$:

$$\begin{aligned}\frac{dZ}{d\alpha} &= - \sum_{i=1}^m D(i) y_i h(x_i) e^{-y_i \alpha_i h(x_i)} = 0 \\ - \sum_{i: y_i = h(x_i)} D(i) e^{-\alpha} + \sum_{i: y_i \neq h(x_i)} D(i) e^{\alpha} &= 0 \\ -e^{-\alpha}(1 - \epsilon) + e^{\alpha}\epsilon &= 0 \\ \alpha &= \frac{1}{2} \log \frac{1 - \epsilon}{\epsilon}\end{aligned}$$

\Rightarrow The minimisator of the upper bound is $\alpha_t = \frac{1}{2} \log \frac{1 - \epsilon_t}{\epsilon_t}$



Boosting Variations and Applications



Variations of Boosting

- **Logitboost, Gentleboost, Floatboost, K1boost,...**
 - How to update the weights
 - How to combine the weak classifiers
- **Example: Gentleboost**
 - Update: $f_m(\mathbf{x}) = P(y=1 \mid \mathbf{x}) - P(y=0 \mid \mathbf{x})$
Instead of likelihood ratio

$$f_m(x) = \frac{1}{2} \log \frac{P_w(y=1|x)}{P_w(y=-1|x)}$$

KLBoost

- **Boost using Kullbeck-Leibler distance**

- Distance between two pdf

$$D(P \parallel Q) = \sum_x P(x) \log \left(\frac{P(x)}{Q(x)} \right)$$

- **Weak classifier**

- Vector(line) projection (inner product)

- **KL Analysis**

- Weak classifier selection
- Feature vector that maximizes KL distance between two classes

KLBoost Procedure (I)

- Given learning set, and initial weights

$$\left\{ (x_1, y_1), (x_2, y_2), \dots, (x_N, y_N), \text{ where } x_i \in \mathbf{R}^d, y_i = \{+1, -1\} \right\} \quad \omega_i^{(0)} = \frac{1}{N}$$

- ❖ Find an optimal feature (projection vector)

$$\phi_k^* = \arg \max_{\phi} KL(\phi) = \int \left[h_k^+(\phi^T x) - h_k^-(\phi^T x) \right] \log \frac{h_k^+(\phi^T x)}{h_k^-(\phi^T x)} d\phi^T x$$

- ❖ Combine weak classifiers

$$\lambda_i(\phi_i^T x) = \alpha_i \log \left(\frac{h_i^+(\phi_i^T x)}{h_i^-(\phi_i^T x)} \right)$$

- ❖ Update the weights

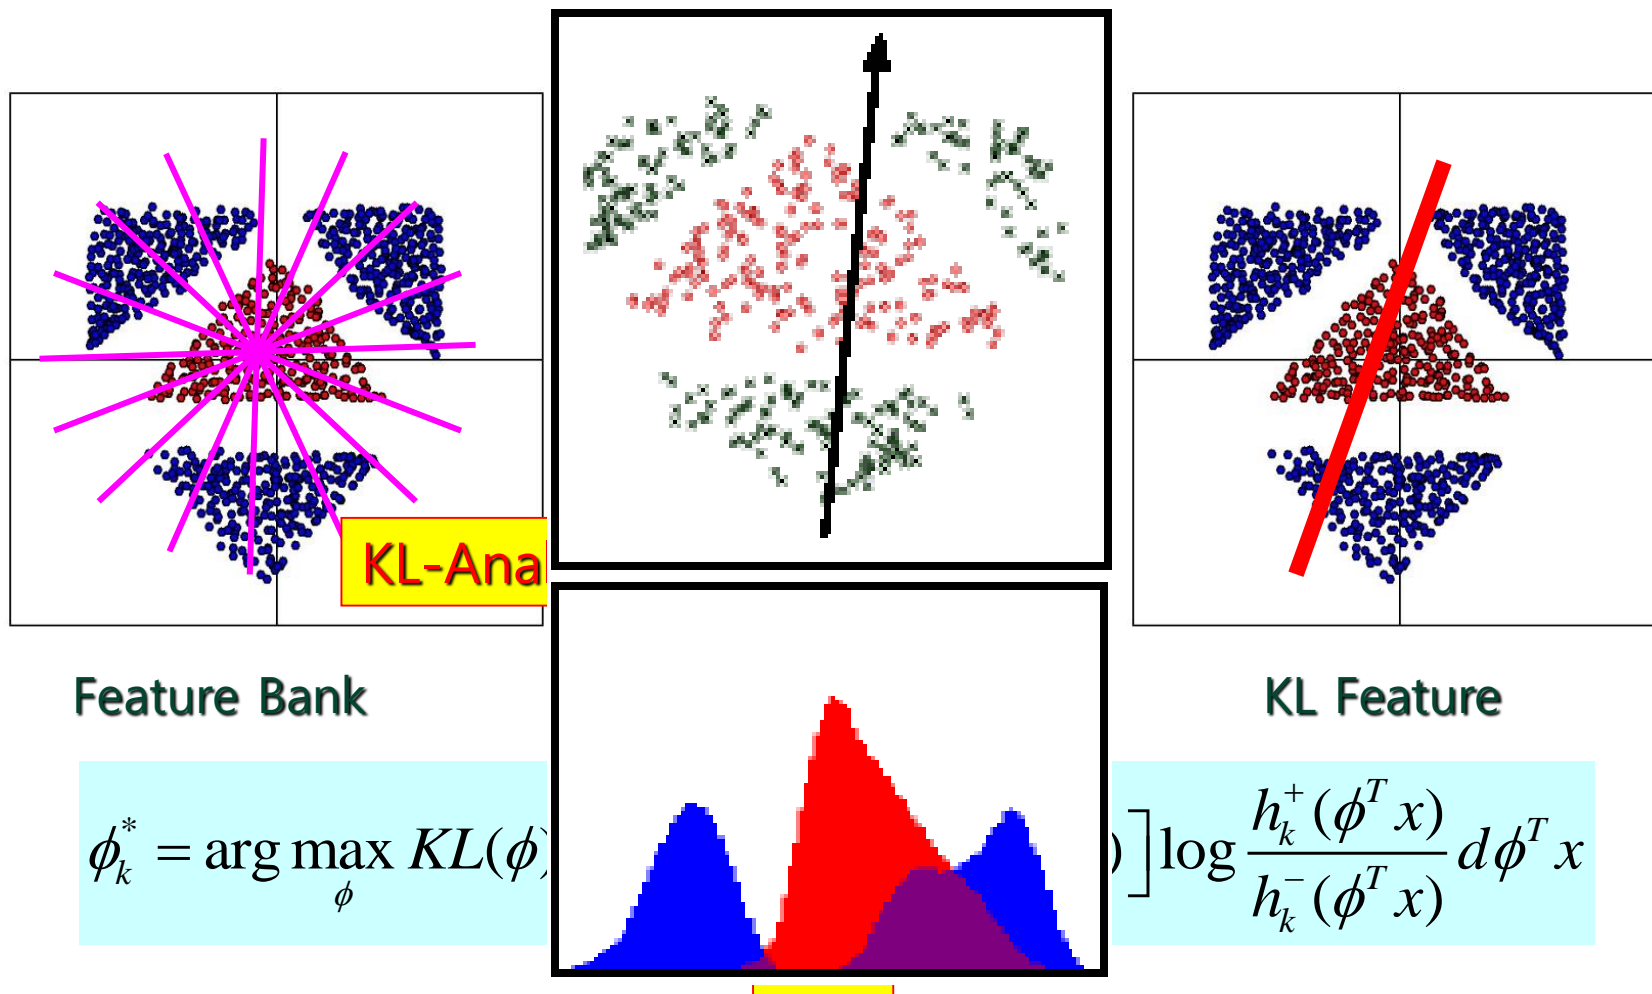
$$\omega_{k+1}(x_i) = \frac{1}{Z_k} \omega_k(x_i) \exp \left\{ -\beta_k y_i F_k(x_i) \right\} \quad \beta \leftarrow \ln \left(\frac{1 - \varepsilon_M}{\varepsilon_M} \right)$$

- ❖ Strong classifier

$$F(x) = \text{sign} \left[\sum_{i=1}^N \lambda_i(\phi_i^T x) \right]$$


Example of KLBoost (I)

- 2D vector classification
 - Line projection and histogram

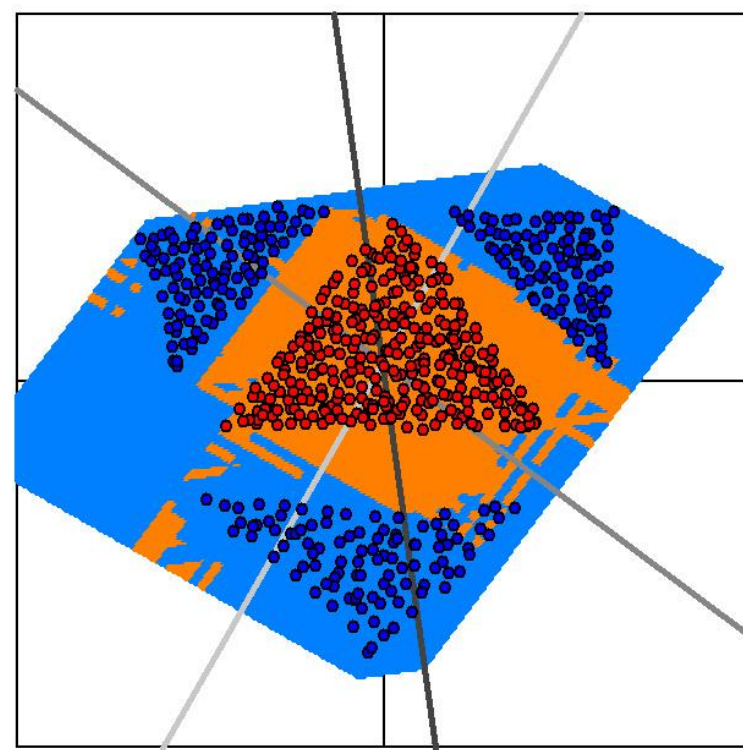


Example of KLBoost (II)

- **KL Boosting example 1**
 - 2D data classification




2-class 2D data




Classified data

Example of KLBoost (III)

- KL Boosting example 2
 - 2D data classification



2-class 2D data



Classified data



Exemplary Application of Boosting


- **For application of boosting process**
 - Design the weak classifiers
 - Design matching schemes
- **Weak classifiers**
 - Easy to generate (weak classifiers bank)
 - Simple calculation
 - Intuitive idea (operation)
- **Example: Car model recognition**
 - What features are used?
 - How are the weak classifiers designed?

2D Haar-Like Patterns


❖ Design some 2D Haar-Like Patterns



Car models



Cascaded tree structure




2D Haar-like Patterns
(features)

2D Haar-Like Patterns




Car models



Car models

Mean outputs

- 120
- 80
- 70
- 65



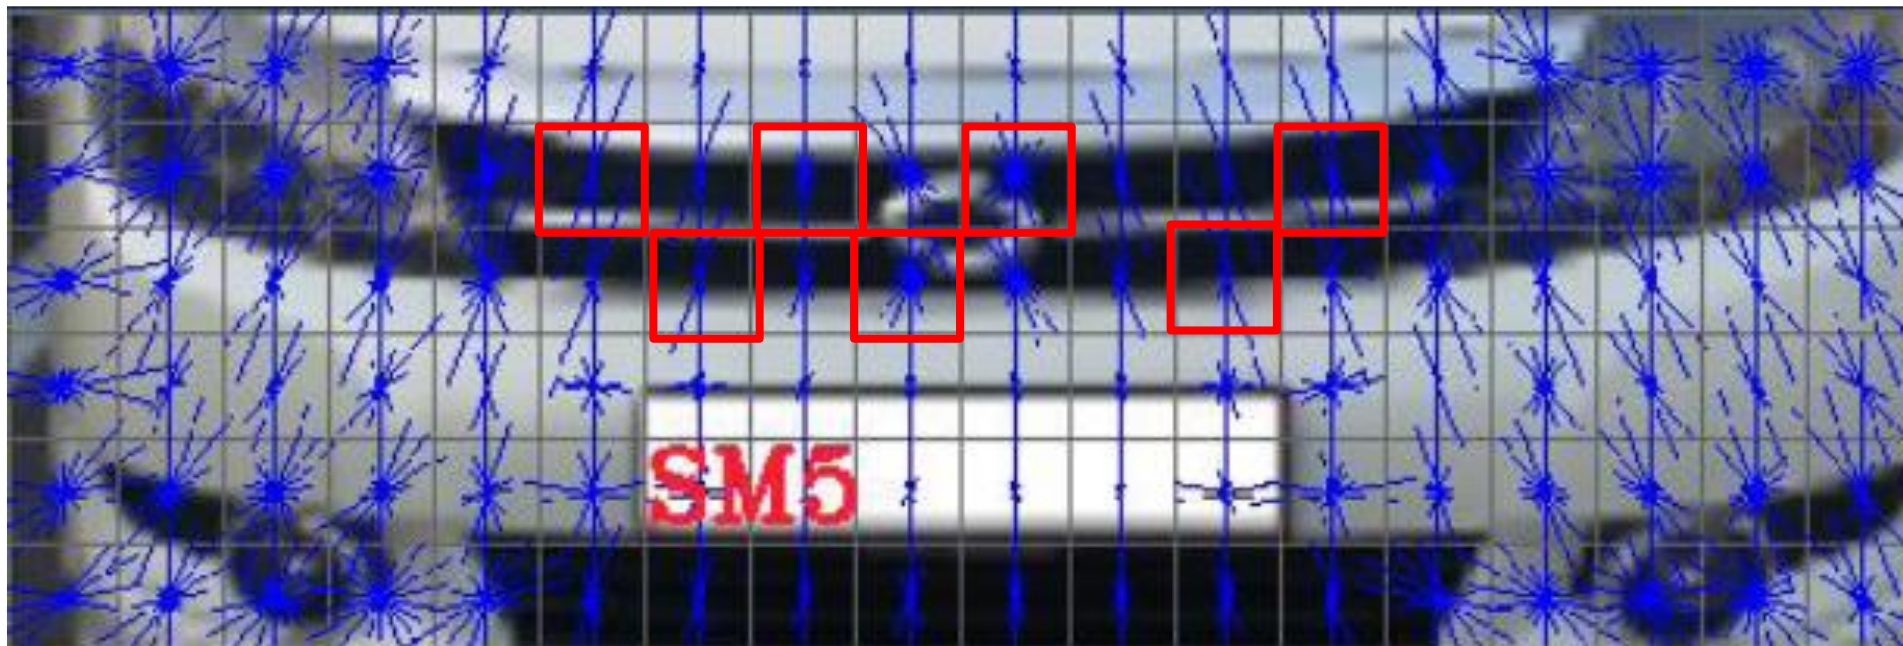
2D Harr-like Patterns
(features)

Mean outputs


- 30
- 40
- 45
- 50

❖ Histogram of Gradient orientations

- Partial distribution
- Select maximally discriminative parts



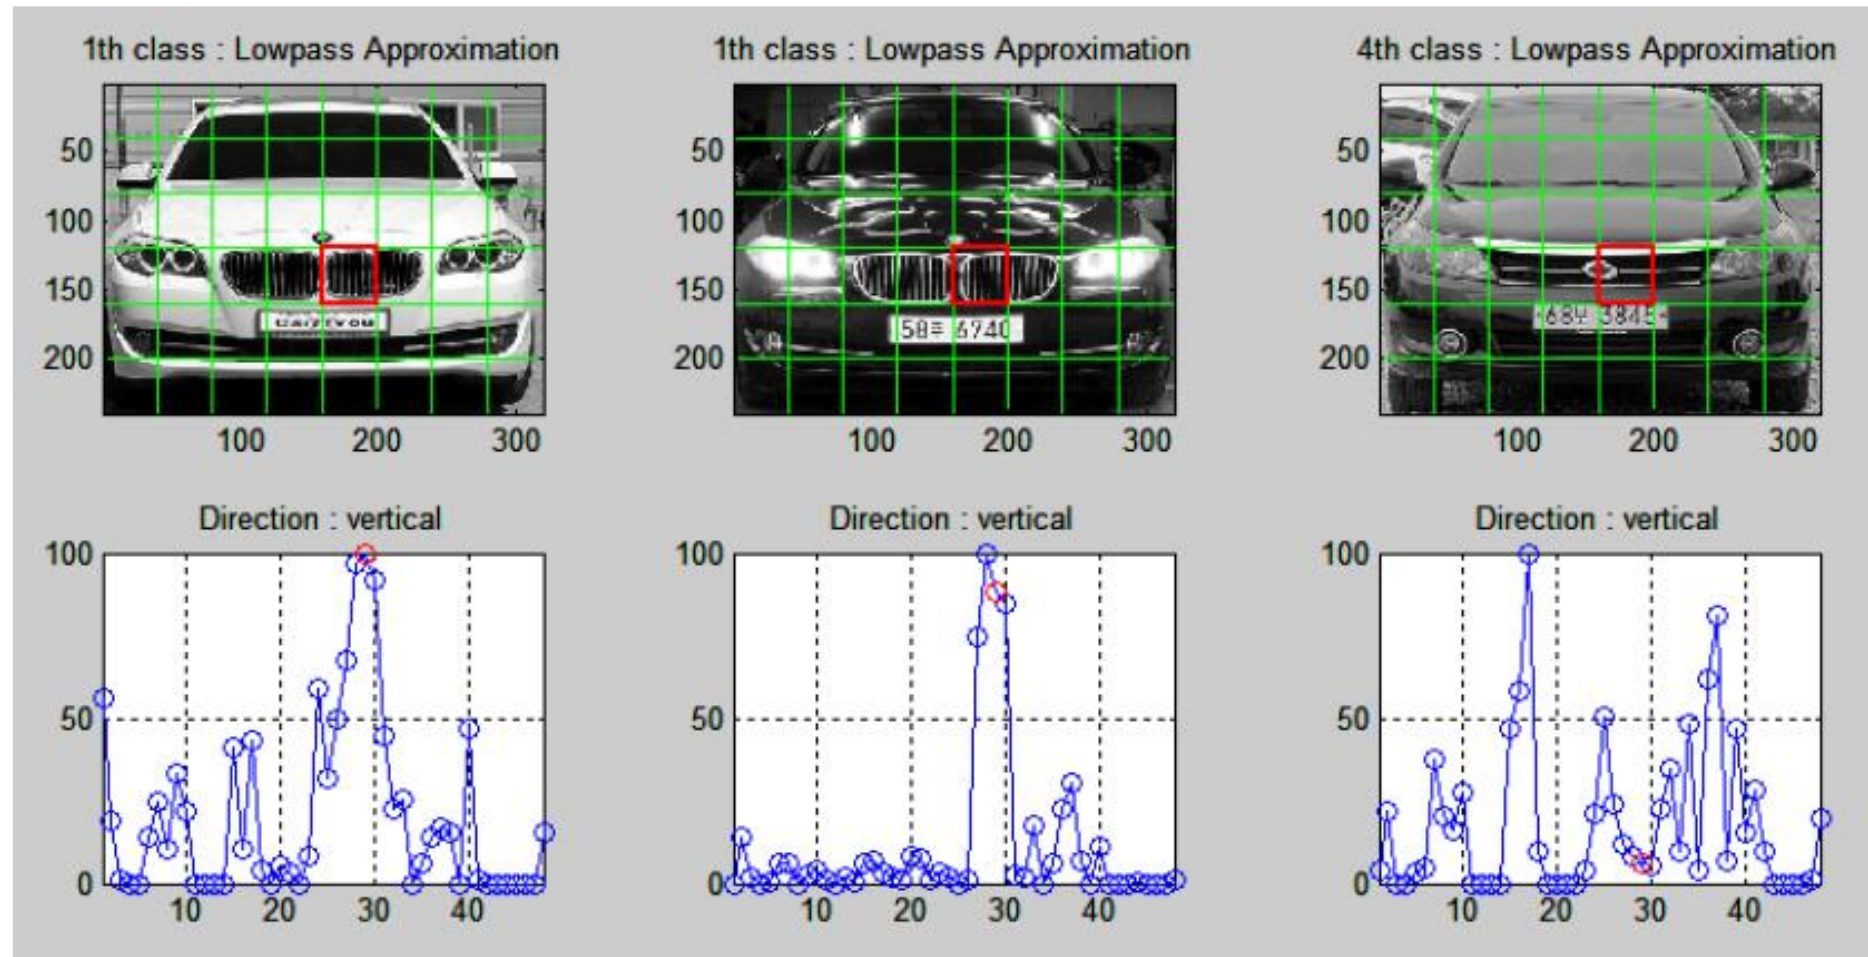
Features: HOG for subblocks (parts)




Most discriminative parts

DFT Coefficients

❖ Specific distribution of DFT Coefficients




Features: Distribution of DFT Coefficients



Most discriminative parts

2D Gabor Patterns

- **Gabor wavelet filter**
 - Texture analysis
 - 5 scales, 8 orientations



2D Gabor Patterns

- Gabor filter response

